



KTH Computer Science
and Communication

Exam in DD2431 Machine Learning 2016-10-28, kl 14.00 – 18.00

Aids allowed: *calculator, language dictionary.*

In order to pass, you have to fulfill the requirements defined both in the A- and in the B-section.

A Questions on essential concepts

Note: As a prerequisite for passing you must give the correct answer to almost *all* questions. Only *one* error will be accepted, so pay good attention here.

A-1 Probabilistic Learning

In maximum likelihood estimation we adjust the model parameters so that:

- a) The probability of the data given the model is maximized.
- b) The training data and the test data have similar distributions.
- c) The classes are linearly separable.

A-2 Regression and Classification

Choose the most proper statement reflecting the output formats of regression and classification.

- a) They are both discrete.
- b) Discrete for classification and continuous-valued for regression.
- c) They are both continuous-valued.

A-3 Shannon Entropy

Consider a single toss of fair coin. Regarding the uncertainty of the outcome {head, tail}:

- a) The entropy is zero.
- b) The entropy is equal to one bit.
- c) The entropy has nothing to do with the measure of uncertainty.

A-4 Overfitting

You have trained a model (classifier) using some training sample data. Under which conditions is overfitting most likely to occur?

- a) Relatively complex model is used where few training samples are available.
- b) Relatively many training samples are used.
- c) Relatively simple model is used.

A-5 Artificial Neural Networks

What happens during *learning* in an artificial neural network?

- a) Weights between nodes are adjusted
- b) Nodes are added and removed
- c) The distribution of training samples is modified

A-6 Support Vector Machine

In a *support vector machine*, what does it mean when a data point has an *alpha*-value of zero?

- a) It is impossible to find a solution with the selected kernel
- b) The point should not be trusted due to its high variance
- c) The point is not used for defining the separating surface

A-7 Ensemble Learning

Which one below best describes the characteristics of Ensemble methods in machine learning?

- a) Ensemble learning is not well-suited to parallel computing.
- b) The priority is to ensure that the models to be combined perform similarly to each other.
- c) Weak learners are trained and combined.

A-8 The Subspace Method

For the subspace methods, a technique of dimensionality reduction is often used to represent the data distribution in each class. Which of these techniques is most suited for this purpose?

- a) Principal Component Analysis (PCA).
- b) Pulse-Code Modulation (PCM).
- c) Phase Change Memory (PCM).

Note: Your answers need be on a solution sheet (**we will not receive this page**).

B Graded problems

A pass is guaranteed with the required points for 'E' below (excluding bonus) in this section *and* the prerequisite in the A-section.

Preliminary number of points required for different grades:

$$22 \leq p \leq 24 \rightarrow A$$

$$19 \leq p < 22 \rightarrow B$$

$$16 \leq p < 19 \rightarrow C$$

$$13 \leq p < 16 \rightarrow D$$

$$8 \leq p < 13 \rightarrow E$$

$$0 \leq p < 8 \rightarrow F$$

B-1 Terminology

(4p)

For each term (a–h) in the left list, find the explanation from the right list which best describes how the term is used in machine learning.

- | | |
|-----------------------------|--|
| | 1) Method for estimating the mean of k observations |
| | 2) Issues in data sparsity in space |
| | 3) Convex optimization |
| | 4) Clustering method based on centroids |
| a) k -means | 5) Probability before observation |
| b) A priori probability | 6) Problems in increasing computation cost |
| c) Curse of dimensionality | 7) The number of base vectors |
| d) Negative sample | 8) Training samples that are considered as outliers |
| e) RANSAC | 9) Algorithm to learn with latent variables |
| f) Error backpropagation | 10) Algorithm to estimate errors |
| g) Dimension of a subspace | 11) The number of elements in a vector |
| h) Expectation Maximization | 12) An approach to train artificial neural networks |
| | 13) Random strategy for amplitude compensation |
| | 14) Training data which is not part of the concept being learned |
| | 15) Robust method to fit a model to data with outliers |
| | 16) Probability at an earlier time |




Figure 1. Illustration for Problem B-2

B-2 Probability based learning

(3p)

In Figure 1 you see two examples of classification problems where the marker \times indicates points from one class and the marker \circ , points from the other class. We call the problem in the *left* plot Problem A and the one on the *right* Problem B. We want to use a Maximum Likelihood classifier based on Gaussian distributions to solve each of these problems. Each class is described by a distribution such that the likelihood of any point x given the class c is:

$$\mathcal{L}(x|c) = \mathcal{N}(x|\mu_c, \Sigma_c) = \frac{1}{2\pi\sqrt{\det(\Sigma_c)}} \exp\left(-\frac{1}{2}(x - \mu_c)^T \Sigma_c^{-1}(x - \mu_c)\right) \quad \text{with } c \in \{\times, \circ\}.$$

- For each Problem A and B: what should the mean vectors μ_c be for each class if we consider a generative approach?
- Consider Problem A. Will the classifier work without errors if we use the mean vectors from question **a** and the identity matrix as covariance for both classes? That is:

$$\Sigma_{\times} = \Sigma_{\circ} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

Hint: you do not need to perform complex derivations to answer this question, if you have a convincing theoretical argument.

- Consider now Problem B. Can you use the same covariance matrices as in question **b** to solve this problem?

Motivate your answers!

NOTE: if you want to refer to the figure, **do not write on this sheet of paper** because we will not receive it with your solutions.




Figure 2. Illustration for Problem B-3

B-3 Probability based Learning: Latent variables

(3p)

Figure 2 shows a set of six data points with coordinates $-5, -4, -3, 3, 4, 5$ indicated by the \times markers. The figure also shows two probability distribution functions d_1 (thick continuous line) and d_2 (thick dashed line). The PDFs have mean $\mu_1 = -1$ and $\mu_2 = 1$ respectively. They have a maximum value of 0.2 corresponding to the mean, and decrease linearly until they reach zero 5 units away from the mean in both directions.

We do not know which distribution has generated each data point, but we assume the two distribution functions are equally likely *a priori*. We want to use Expectation Maximization to obtain a new estimate of the mean parameter for each distribution, given the data. To do this, perform the following steps:

- with the help of the plot, evaluate each probability distribution function on each data point, that is, compute $p(x_i|d_j)$. **Hint:** define a symbol $\alpha = \frac{0.2}{5}$ to simplify the calculations throughout this exercise.
- perform one Expectation step, that is, evaluate the responsibilities $p(d_j|x_i)$ for each point and each distribution.
- perform one Maximization step, that is, compute the new values for the means of the two distributions. **Hint:** for a set of data points x_i and a set of generic weights w_i , the general formula for the weighted average is:

$$a = \frac{\sum_i w_i x_i}{\sum_i w_i}$$

B-4 Classification

(3p)

Suppose that we take a data set, divide it into equally-sized training and test sets, and then try out two different classification procedures. First we use the Subspace Method and get an error rate of 12% on the training data. We also get the average error rate (averaged over both test and training data sets) of 20%. Next we use k -nearest neighbor (where $k = 1$) and get an average error rate (averaged over both test and training data sets) of 15%.

- a) What was the error rate with the Subspace Method on the test set?
- b) Based on these results, indicate the method which we should prefer to use for classification of new observations, with a simple reasoning.
- c) Now, if the result of k -nearest neighbor (average error rate being 15%) was based on $k = 5$, would we know if the same conclusion drawn in **b** should still hold? Answer *yes* or *no*, with a short reasoning.

B-5 Decision Forests

(2p)

Choose the correct answers in the following questions on Decision Forests.

- a) Mainly two kinds of randomness are known to form the principle of Decision Forests. In which two of the following processes are those randomnesses involved?
 - i. In deciding the number of trees used.
 - ii. In the rule of terminaing a node as a leaf node.
 - iii. In the way to formulate the information gain.
 - iv. In generating bootstrap replicas.
 - v. In feature selection at each node.
- b) Suppose we have generated a Decision Forest using five bootstrapped samples from a data set containing three classes, {Green, Blue, Red}. We then applied the forest to a specific test input, x , and observed five estimates of $P(\text{Class is Green}|x)$: 0.3, 0.4, 0.55, 0.6, and 0.75.
There are two common ways to combine these results together into a single class prediction. One is the majority vote approach, and the other is based on the average probability. In this example, what is the final classification under each of these two approaches? Motivate your answer by short phrases.
 - i. Green in majority vote and Red in averaging.
 - ii. Blue in majority vote and Green in averaging.
 - iii. Green in both approaches.
 - iv. Blue or Red in both approaches.
 - v. Blue or Red in majority vote and Green in averaging.

B-6 Regression with regularization: LASSO

(3p)

For a set of N training samples $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$, each consisting of input vector \mathbf{x} and output y , suppose we estimate the regression coefficients $\mathbf{w}^\top (\in \mathbf{R}^d) = \{w_1, \dots, w_d\}$ in a linear regression model by minimizing

$$\sum_{n=1}^N (y_n - \mathbf{w}^\top \mathbf{x}_n)^2 \quad \text{subject to} \quad \sum_{i=1}^d |w_i| \leq s$$

for a particular value of s .

For parts **a)** through **c)**, indicate which of i. through v. is correct. Briefly justify your answer.

- a)** As we increase s from 0, the *training* error (residual sum of squares, RSS) will:
- i. Steadily increase.
 - ii. Steadily decrease.
 - iii. Remain constant.
 - iv. Increase initially, and then eventually start decreasing in an inverted U shape.
 - v. Decrease initially, and then eventually start increasing in a U shape.
- b)** Repeat **a)** for *test* RSS.
- c)** Repeat **a)** for variance and (squared) bias, respectively.

B-7 Support Vector Machines

(3p)

Training a support vector machine using a quadratic kernel

$$\mathcal{K}(\vec{x}, \vec{y}) = (\vec{x}^T \vec{y} + 1)^2$$

has resulted in the following three support vectors:

$$s_1 = \begin{bmatrix} 0 \\ 0 \end{bmatrix} \quad s_2 = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \quad s_3 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

The first (s_1) is a negative sample with $\alpha = \frac{7}{3}$ while the other two are positive samples where $\alpha = \frac{2}{3}$.

a) Determine how the following *new* datapoints will be classified:

$$x_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \quad x_2 = \begin{bmatrix} 2 \\ 0 \end{bmatrix} \quad x_3 = \begin{bmatrix} -2 \\ 0 \end{bmatrix} \quad x_4 = \begin{bmatrix} -3 \\ 0 \end{bmatrix}$$

You must show the formulas and calculations used to arrive at your answer.

b) Draw a diagram of the input space with all the relevant points ($s_1, s_2, s_3, x_1, x_2, x_3, x_4$) and show the shape and position of the decision boundary.

B-8 Perceptron Learning

(3p)

Consider the training data in the table, where + means a positive sample and – a negative. Show, step by step, the values of the the weight vector as it is updated when these samples are used to train a linear classifier using the *perceptron learning rule*. Use the samples in the order as they are listed, and repeat them until all samples are correctly classified.

To make the calculations simple; start by setting all weights to zero, and use a step size (η) of 1. Further, use a threshold function which gives a negative (low) output when the input happens to be *exactly on* the threshold.

Do not only list the successive values of the weights, but also explain what calculations are being done.

1	0	1	0	0	+
1	1	1	1	0	–
0	1	0	1	0	+
1	1	1	0	1	–
0	0	1	0	1	+
1	1	0	1	1	–