




Artificial Neural Networks





Feed Forward Networks

Applications

Classical Examples

Multi Layer Networks

Possible Mappings

Backprop Algorithm

Practical Problems

Deep Networks

Vanishing Gradients

Convolutional Networks



Feed Forward Networks

Applications

Classical Examples

Multi Layer Networks

Possible Mappings

Backprop Algorithm

Practical Problems

Deep Networks

Vanishing Gradients


Convolutional Networks

Artificial Neural Networks (ANN)

- ▶ Inspired from the nervous system
- ▶ Parallel processing

We will focus on **one** class of ANNs:

Feed-forward Layered Networks



Applications

Operates like a general "Learning Box"!


Classification



Function Approximation



Multidimensional Mapping



ALVINN


Autonomous driving



Trained to mimic the behavior of human drivers

NetTalk

Speech Synthesis



Trained using a large database of spoken text



Feed Forward Networks

Applications

Classical Examples

Multi Layer Networks

Possible Mappings


Backprop Algorithm

Practical Problems

Deep Networks

Vanishing Gradients

Convolutional Networks



A two layer network can implement arbitrary decision surfaces
...provided we have *enough hidden units*

How can we train a multi layer network?



Neither perceptron learning, nor the delta rule can be used


Fundamental problem:

When the network gives the wrong answer
there is no information on in which direction
the weights need to change to improve the result


Trick:

Use threshold-like, but **continuous** functions






First layer response



Sum of base functions



Learning strategy:

Minimize the error (E) as a function of **all** weights (\vec{w})

1. Compute the direction in weight space where the error increases the most $\text{grad}_{\vec{w}}(E)$
2. Change the weights in the opposite direction


$$w_i \leftarrow w_i - \eta \frac{\partial E}{\partial w_i}$$

Normally one can use the error from each example separately

$$E = \frac{1}{2} \sum_{k \in \text{Out}} (t_k - y_k)^2$$

A common "threshold-like function" is

$$\rho(x) = \frac{1}{1 + e^{-x}}$$



The gradient can be expressed as a function of a *local generalized error* δ

$$\frac{\partial E}{\partial w_{ji}} = -\delta_i x_j \quad w_{ji} \leftarrow w_{ji} + \eta \delta_i x_j$$

Output layer:

$$\delta_k = y_k \cdot (1 - y_k) \cdot (t_k - y_k)$$

Hidden layers:

$$\delta_j = h_j \cdot (1 - h_j) \cdot \sum_k w_{kj} \delta_k$$

The errors δ propagate backwards through the layers
Error backpropagation (BackProp)

Things to think about when using BackProp

- ▶ Sloooooow
Normal to require thousands of iterations through the dataset
- ▶ Gradient following
Risk of getting stuck in local minima
- ▶ Many parameters
 - ▶ Step size η
 - ▶ Number of layers
 - ▶ Number of hidden units
 - ▶ Input and output representation
 - ▶ Initial weights



Feed Forward Networks

Applications

Classical Examples

Multi Layer Networks

Possible Mappings


Backprop Algorithm

Practical Problems

Deep Networks

Vanishing Gradients

Convolutional Networks




Deep networks — Networks with many layers

- ▶ Error gradients become smaller from layer to layer
- ▶ Pure Backprop becomes unusable for deep networks

Deep Belief Networks

- ▶ Unsupervised learning of features
Restricted Boltzmann Machine
- ▶ Greedy learning from the bottom, layer by layer
Optimize for ability to reconstruct previous layer
- ▶ Supervised Backprop to finalize classifier

Convolutional Networks



- ▶ Alternating convolution and subsampling layers
- ▶ Weight sharing
- ▶ Trained using Backprop